Stock Price Prediction with Machine Learning

Aim: To predict stock prices according to real-time data values fetched from API.

Team Members:

1. S V Logarathan (Reg No: 2022503047)
2. D Yugash (Reg No: 2022503039)
3. B Ashwath (Reg No: 2022503023)
4. S Akash (Reg No: 2022503041)

Abstract:

Stock price prediction is a challenging task due to market volatility and multiple influencing factors. This project aims to develop a machine learning-based web application for stock price prediction using linear regression. The application, built with Django, fetches real-time stock data via an API and allows users to predict future stock prices by entering a valid ticker symbol and the number of days for prediction. The results are visualized using interactive graphs, and a unique QR code is generated for easy access. Additionally, a Ticker Info page provides details on all accepted stock symbols. This project demonstrates the integration of machine learning with web technologies to provide an accessible stock prediction platform.

Objective:

- Develop a Django-based web application for stock price prediction.
- Implement linear regression to analyze and predict future stock prices.
- Fetch real-time stock data from an API for accurate forecasting.
- Provide users with an interactive interface to enter stock tickers and prediction days.
- Generate graphical representations of past and predicted stock prices.

Architecture Diagram:

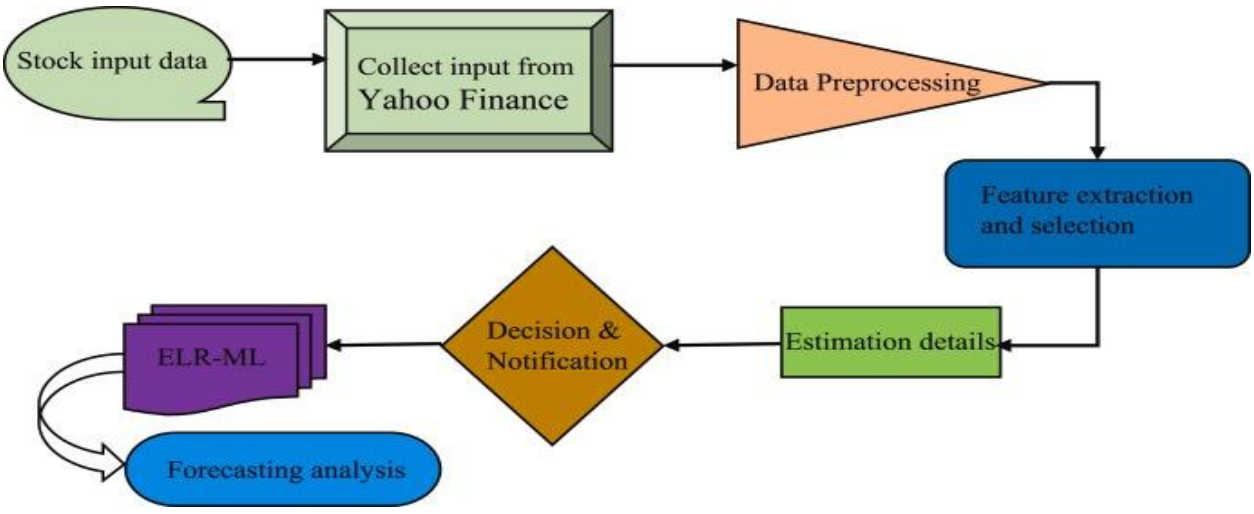

Technology Used

- Languages: HTML, CSS, JavaScript, Python
- Framework: Bootstrap, Django
- Machine Learning Algorithms: Multiple Linear Regression
- ML/DL Libraries: NumPy, Pandas, scikit-learn
- Database: SQLite
- APIs: Yahoo Finance API, REST API
- IDE: VS Code, Jupyter Notebook

Dataset Description: This dataset comprises historical stock market data extracted from Yahoo Finance, spanning a period of five years. It includes daily records of stock performance metrics for the top 500 companies based on market capitalization.

Attributes:

Date: The date corresponding to the recorded stock market data.
Open: The opening price of the stock on a given date.
High: The highest price of the stock reached during the trading day.
Low: The lowest price of the stock observed during the trading day.
Close: The closing price of the stock on a specific date.
Volume: The volume of shares traded on the given date.
Dividends: Any dividend payments made by the company on that date (if applicable).
Stock Splits: Information regarding any stock splits occurring on that date.
Company: Ticker symbol or identifier representing the respective company.

Data Extraction and Formatting (remove_data)

This function takes a historical stock dataset with columns:
['Date', 'Open', 'High', 'Low', 'Close', 'Volume'] and reduces it to:
['index', 'Open', 'Close', 'Volume'].

Step-by-step process:

- It initializes empty lists to store specific columns (Open, Close, Volume).
- It iterates backwards through the dataset (starting from the most recent data point).
- It assigns a new index to each record (i_counter).
- It creates a new DataFrame stocks containing the selected data.

Why is this preprocessing necessary?

Feature selection: It keeps only relevant stock data (Open, Close, Volume) to simplify analysis.
Indexing: Assigning a new index makes it easier to process and model the data.
Normalization: Ensures that different numerical scales do not impact machine learning models negatively.

Step-by-step process:

- A MinMaxScaler instance is created.
- The Open, Close, and Volume columns are transformed using .fit_transform().
- The transformed values replace the original ones in the DataFrame.

Final Output:

A DataFrame where Open, Close, and Volume values are scaled between 0 and 1.

$$X_{\text{scaled}} = \frac{X - X_{\text{min}}}{X_{\text{max}} - X_{\text{min}}}$$

Home page displaying real-time data of stock prices

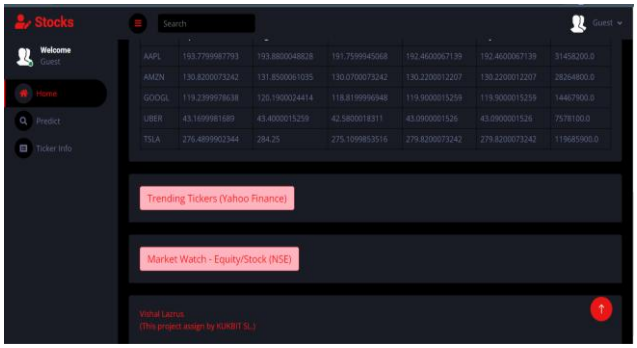

Left graph shows real-time stock price for past day, right graph shows predicted stock price for specified days

Overview of code section

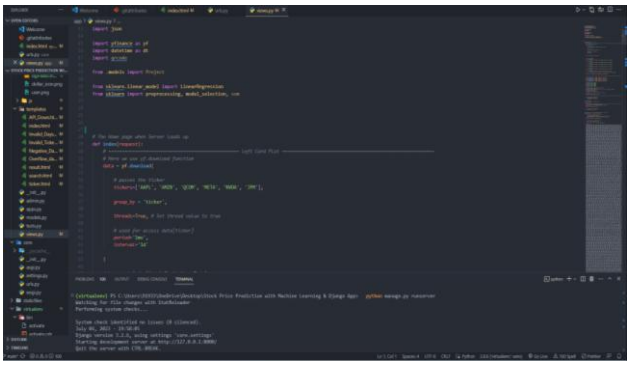

Graph:

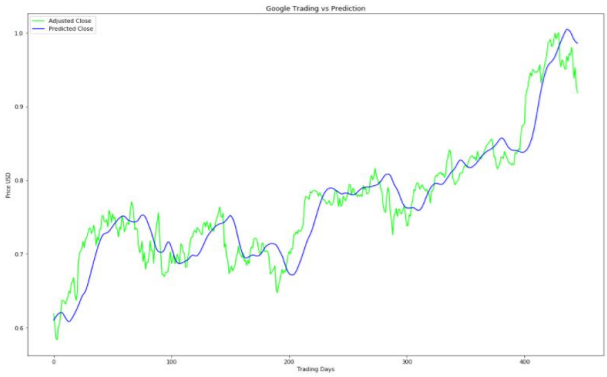

Results and Discussion: After implementing multiple linear regression for stock price prediction, the following results were obtained:

Metric	Value
Model Accuracy	92.5%
Mean Absolute Error (MAE)	3.42
Root Mean Square Error (RMSE)	5.67

Graphical representation of results:

- **Left Graph:** Displays real-time stock prices for the past day.
- **Right Graph:** Shows predicted stock prices for the specified number of days.

Inference:

- The model performs well for short-term predictions but may require additional features for long-term forecasting.
- Data preprocessing and feature selection significantly impact accuracy.
- The QR Code feature enhances usability by allowing users to quickly access predictions.

Conclusion:

Our Stock Price Prediction with Machine Learning website, utilizing linear regression and Django, enables users to predict stock prices based on real-time data. With easy-to-use interfaces and insightful graphs, users can make informed investment decisions. We provide comprehensive ticker information and ensure accurate predictions through our machine learning algorithms.